

WEEK-06

Average temperature, Mean Max temperature

```
hduser@bmsce-Precision-T1700:~$ hadoop jar /home/hduser/Desktop/neannxtemp.jar MeanMaxDriver /yathri/weather1.txt outputtempmax
23/06/10 10:03:53 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
23/06/10 10:03:53 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
23/06/10 10:03:53 WARN mapreduce.JobSubmitter: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
23/06/10 10:03:53 INFO input.FileInputFormat: Total input paths to process : 1
23/06/10 10:03:53 INFO mapreduce.JobSubmitter: number of splits:1
23/06/10 10:03:53 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local86685270_0001
23/06/10 10:03:53 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
23/06/10 10:03:53 INFO mapreduce.Job: Running job: job_local86685270_0001
23/06/10 10:03:53 INFO mapred.LocalJobRunner: OutputCommitter set in config null
23/06/10 10:03:53 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
23/06/10 10:03:53 INFO mapred.LocalJobRunner: Waiting for map tasks
23/06/10 10:03:53 INFO mapred.LocalJobRunner: Starting task: attempt_local86685270_0001_m_000000_0
23/06/10 10:03:53 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
23/06/10 10:03:53 INFO mapred.MapTask: Processing split: hdfs://localhost:54310/yathri/weather1.txt:0+888190
23/06/10 10:03:53 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
23/06/10 10:03:53 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
23/06/10 10:03:53 INFO mapred.MapTask: soft limit at 83886080
23/06/10 10:03:53 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
23/06/10 10:03:53 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
23/06/10 10:03:53 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
23/06/10 10:03:53 INFO mapred.LocalJobRunner:
23/06/10 10:03:53 INFO mapred.MapTask: Starting flush of map output
23/06/10 10:03:53 INFO mapred.MapTask: Spilling map output
23/06/10 10:03:53 INFO mapred.MapTask: bufstart = 0; bufend = 45948; bufvoid = 104857600
23/06/10 10:03:53 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26188144(104752576); length = 26253/6553600
23/06/10 10:03:53 INFO mapred.MapTask: Finished spill 0
23/06/10 10:03:53 INFO mapred.Task: Task:attempt_local86685270_0001_m_000000_0 is done. And is in the process of committing
23/06/10 10:03:53 INFO mapred.LocalJobRunner: map
23/06/10 10:03:53 INFO mapred.Task: Task 'attempt_local86685270_0001_m_000000_0' done.
23/06/10 10:03:53 INFO mapred.LocalJobRunner: Finishing task: attempt_local86685270_0001_m_000000_0
```

```
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls outputtempmax
Found 2 items
-rw-r--r-- 1 hduser supergroup 0 2023-06-10 10:03 outputtempmax/_SUCCESS
-rw-r--r-- 1 hduser supergroup 74 2023-06-10 10:03 outputtempmax/part-r-00000
```

Bytes Written=72

```
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls outputtempmax1
Found 2 items
-rw-r--r-- 1 hduser supergroup 0 2023-06-10 10:07 outputtempmax1/_SUCCESS
-rw-r--r-- 1 hduser supergroup 72 2023-06-10 10:07 outputtempmax1/part-r-00000
hduser@bmsce-Precision-T1700:~$ hadoop fs -cat outputtempmax1/part-r-00000
```

```
01 4
02 1
03 4
04 24
05 78
06 119
07 145
08 146
09 104
10 45
11 23
12 4
```

```
hduser@bmsce-Precision-T1700:~$ hadoop fs -cat outputtempmax/part-r-00000
```

```
01      4
02      0
03      7
04     44
05    100
06    168
07    219
08    198
09    141
10    100
11     19
12      3
```

```
hadoop@bmscece-HP-Elite-Tower-600-G9-Desktop-PC: $ hdfs dfs -copyFromLocal /home/hadoop/Desktop/weather.txt /yathri
hadoop@bmscece-HP-Elite-Tower-600-G9-Desktop-PC: $ hadoop fs -ls /yathri
Found 3 items
drwxr-xr-x 1 hadoop supergroup 0 2023-05-17 09:33 /yathri/Desktop
-rw-r--r-- 1 hadoop supergroup 97 2023-05-17 09:35 /yathri/wc.txt
-rw-r--r-- 1 hadoop supergroup 888978 2023-05-17 10:30 /yathri/weather.txt
hadoop@bmscece-HP-Elite-Tower-600-G9-Desktop-PC: $ hadoop jar /home/hadoop/Documents/jar/Weather.jar AverageDriver /yathri/weather.txt /output2
2023-05-17 10:33:02,346 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2023-05-17 10:33:02,380 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2023-05-17 10:33:02,381 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2023-05-17 10:33:02,432 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2023-05-17 10:33:02,465 INFO input.FileInputFormat: Total input files to process : 1
2023-05-17 10:33:02,490 INFO mapreduce.JobSubmitter: number of splits:1
2023-05-17 10:33:02,546 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local212143084_0001
2023-05-17 10:33:02,546 INFO mapreduce.JobSubmitter: Executing with tokens: []
2023-05-17 10:33:02,599 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2023-05-17 10:33:02,599 INFO mapreduce.Job: Running job: job_local212143084_0001
2023-05-17 10:33:02,600 INFO mapred.LocalJobRunner: OutputCommitter set in Config null
2023-05-17 10:33:02,603 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2023-05-17 10:33:02,603 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2023-05-17 10:33:02,603 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2023-05-17 10:33:02,636 INFO mapred.LocalJobRunner: Waiting for map tasks
2023-05-17 10:33:02,636 INFO mapred.LocalJobRunner: Starting task: attempt_local212143084_0001_m_000000_0
2023-05-17 10:33:02,645 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2023-05-17 10:33:02,645 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2023-05-17 10:33:02,651 INFO mapred.Task: Using ResourceCalculatorProcessTree: [ ]
2023-05-17 10:33:02,652 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/yathri/weather.txt:0+888978
2023-05-17 10:33:02,681 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
2023-05-17 10:33:02,681 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2023-05-17 10:33:02,681 INFO mapred.MapTask: soft limit at 83886080
2023-05-17 10:33:02,681 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2023-05-17 10:33:02,681 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2023-05-17 10:33:02,683 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2023-05-17 10:33:02,748 INFO mapred.LocalJobRunner:
2023-05-17 10:33:02,749 INFO mapred.MapTask: Starting flush of map output
2023-05-17 10:33:02,749 INFO mapred.MapTask: Spilling map output
2023-05-17 10:33:02,749 INFO mapred.MapTask: bufstart = 0; bufend = 59085; bufvoid = 104857600
2023-05-17 10:33:02,749 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26188140(104752560); length = 26257/6553600
2023-05-17 10:33:02,756 INFO mapred.MapTask: Finished spill 0
2023-05-17 10:33:02,759 INFO mapred.Task: Task:attempt_local212143084_0001_m_000000_0 is done. And is in the process of committing
2023-05-17 10:33:02,761 INFO mapred.LocalJobRunner: map
2023-05-17 10:33:02,761 INFO mapred.Task: Task 'attempt_local212143084_0001_m_000000_0' done.
2023-05-17 10:33:02,763 INFO mapred.Task: Final Counters for attempt_local212143084_0001_m_000000_0: Counters: 23
File System Counters
FILE: Number of bytes read=4327
FILE: Number of bytes written=713168
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=888978
HDFS: Number of bytes written=0
HDFS: Number of read operations=5
HDFS: Number of large read operations=0
```

Bytes Written=8

```
hadoop@bmscece-HP-Elite-Tower-600-G9-Desktop-PC:~$ hadoop fs -ls /output2
```

```
Found 2 items
```

```
-rw-r--r-- 1 hadoop supergroup 0 2023-05-17 10:33 /output2/_SUCCESS
-rw-r--r-- 1 hadoop supergroup 8 2023-05-17 10:33 /output2/part-r-00000
```

```
hadoop@bmscece-HP-Elite-Tower-600-G9-Desktop-PC:~$ hadoop fs -cat /output2/part-r-00000
```

```
1902    21
```

```
hadoop@bmscece-HP-Elite-Tower-600-G9-Desktop-PC:~$ █
```